

# Parallel Platforms & Programming Models

# Outline

1. Parallel platforms (~hardware) and programming models (~software)
  - Note: Parallel machine may or may not be tightly coupled to programming model
2. Data dependence
  - Reading: Kumar – ch 1; Hwang – ch1; Foster – ch 1

# Flavors of Parallelism

- Data parallelism:
  - Definition?
- Task parallelism:
  - Definition?

# Data and Task Parallelization

- Data parallel:

```
for (i=0;i<1000;i++)  
    a[i]=b[i]+c[i];
```

- Task parallel:

```
for (i=0;i<1000;i++) /*block 1 */  
    b[i+1]=b[i]+c[i]
```

...

```
for (j=0;j<5;j++) /*block 2*/  
    a[j+1]=a[j]+d[j];
```

# Parallel Platforms

- Basic components of any architecture:
  - Processors and memory (processing units)
  - Interconnect network
- Logic classification based on:
  - Control mechanism (Flynn's Taxonomy)
    - SISD (Single Instruction Single Datastream)
    - SIMD (Single Instruction Multiple Datastream)
    - MISD (Multiple Instruction Single Datastream)
    - MIMD (Multiple Instruction Multiple Datastream)
  - Address space organization
    - Shared Address Space
    - Distributed Address Space

# Parallel Platforms based on Flynn's Taxonomy

# Flynn's Taxonomy

<b>SISD</b> Single Instruction Stream Single Data Stream	<b>SIMD</b> Single Instruction Stream Multiple Data Stream
<b>MISD</b> Multiple Instruction Stream Single Data Stream	<b>MIMD</b> Multiple Instruction Stream Multiple Data Stream

Single  
Data

Multiple  
Data

Single  
Instruction

*SISD*  
typical thread

*SIMD*  
vector processors  
GPUs  
SSE instructions

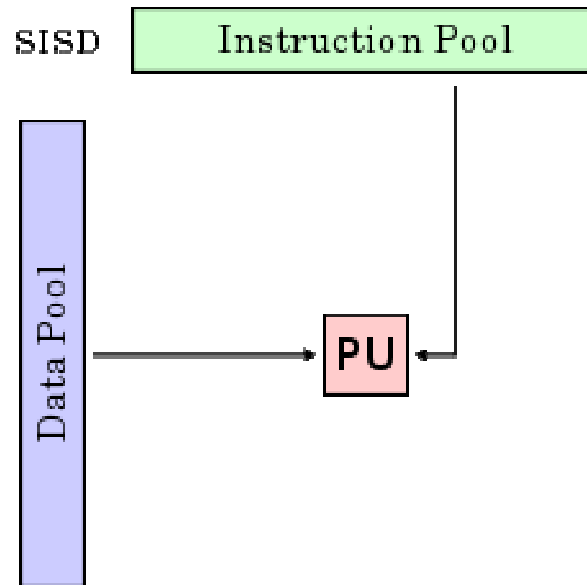
Multiple  
Instruction

*MISD*  
rare  
possibly set of  
filters

*MIMD*  
cluster of computers

# SISD Architecture


- Model of serial Von Neumann machine
- Logically, single control processor
- Includes some supercomputers, such as the 1963 CDC6600 (perhaps the first supercomputer)





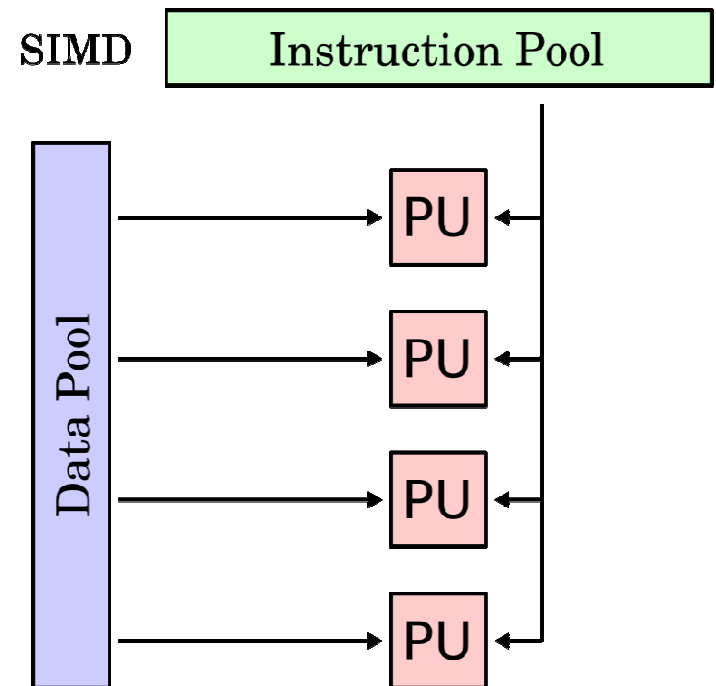
# Von Neumann Architecture

- John von Neumann first authored the general requirements for an electronic computer in 1945
- Aka “stored-program computer”
  - Both program inst. and data are kept in electronic memory
- Since then, all computers have followed this basic design
- Four main components: memory, control unit, ALU, I/O



# SIMD Architecture

- Parallelism achieved by dividing data among the processors
  - Multiple processors execute the same instruction
  - Data that each processor sees may be different
  - Individual processors can be turned on/off at each cycle (“masking”)
- Examples:
  - Many early parallel computers like Illiac IV, Thinking Machines’ CM-2, ...
  - Today, GPU, vector units, and co-processors



# Example of SIMD Vector Units

- Scalar processing
  - Traditional mode
  - One operation produces one result
- SIMD vector units
  - One operation produces multiple results

X

+

Y

X + Y

X

x3 x2 x1 x0

+

Y

y3 y2 y1 y0

X + Y

x3+y3 x2+y2 x1+y1 x0+y0

# SIMD Drawbacks

- Discussion?


# The ill-fated Illiac IV

- Project started in 1965, predicted to cost \$8M and provide 1000 MFLOP/S.
- Delivered to NASA Ames in 1972, cost \$31M, ran first application in 1976, performed 15 MFLOP/S.
- 64 processors, 13-MHz clock, 1MB memory



# Thinking Machine CM2

- CM2 (1990, built by Thinking Machines Corp) had 8,192 to 65,536 one-bit processors, plus one floating-point unit.
- Data Vault provides peripheral mass storage
- Single program - all unmasked operations happened in parallel.



# Vector Processors

- Operate on arrays or vectors of data, while conventional CPU's operate on individual data elements or scalars
- Vector registers
  - Capable of storing a vector of operands and operating simultaneously on their contents
- Vectorized and pipelined functional units
  - The same operation is applied to each element in the vector
- Examples:
  - Cray supercomputers (X-MP, Y-MP, C90, T90, SV1, ...), Fujitsu (VPPxxx), NEC, Hitachi
  - Earth Simulator from Japan (on the TOP500 list)
  - many of these have multiple vector processors, but typically separate processors are used for separate jobs.


# Vector Processors – Pros & Cons

- Discussion?




# Graphic Processing Units (GPU)

- Real time graphics application programming interfaces or API's use points, lines, and triangles to internally represent the surface of an object
- A graphics processing pipeline converts the internal representation into an array of pixels that can be sent to a computer screen
- Several stages of this pipeline (called shader functions) are programmable
  - Typically just a few lines of C code



# MIMD Architecture

- Supports multiple simultaneous instruction streams operating on multiple data streams
  - Each processor executes program independent of other processors
  - Processors operate on separate data streams
- Typically consist of a collection of fully independent PUs or cores, each of which has its own control unit and its own ALU.
- Examples:
  - Current generation systems




# Examples



**Tianhe-2**



**Mira**



**Titan**

# MISD Architectures

- Multiple Instruction Single Data
- Few (if any) actual examples of this class of parallel computer have ever existed.
- The term isn't used (except when discussing the Flynn taxonomy) .
- Perhaps applies to pipelined computation, e.g. sonar data passing through sequence of special-purpose signal processors.


# SIMD vs MIMD

- SIMD platforms
  - ???
- MIMD platforms
  - ???

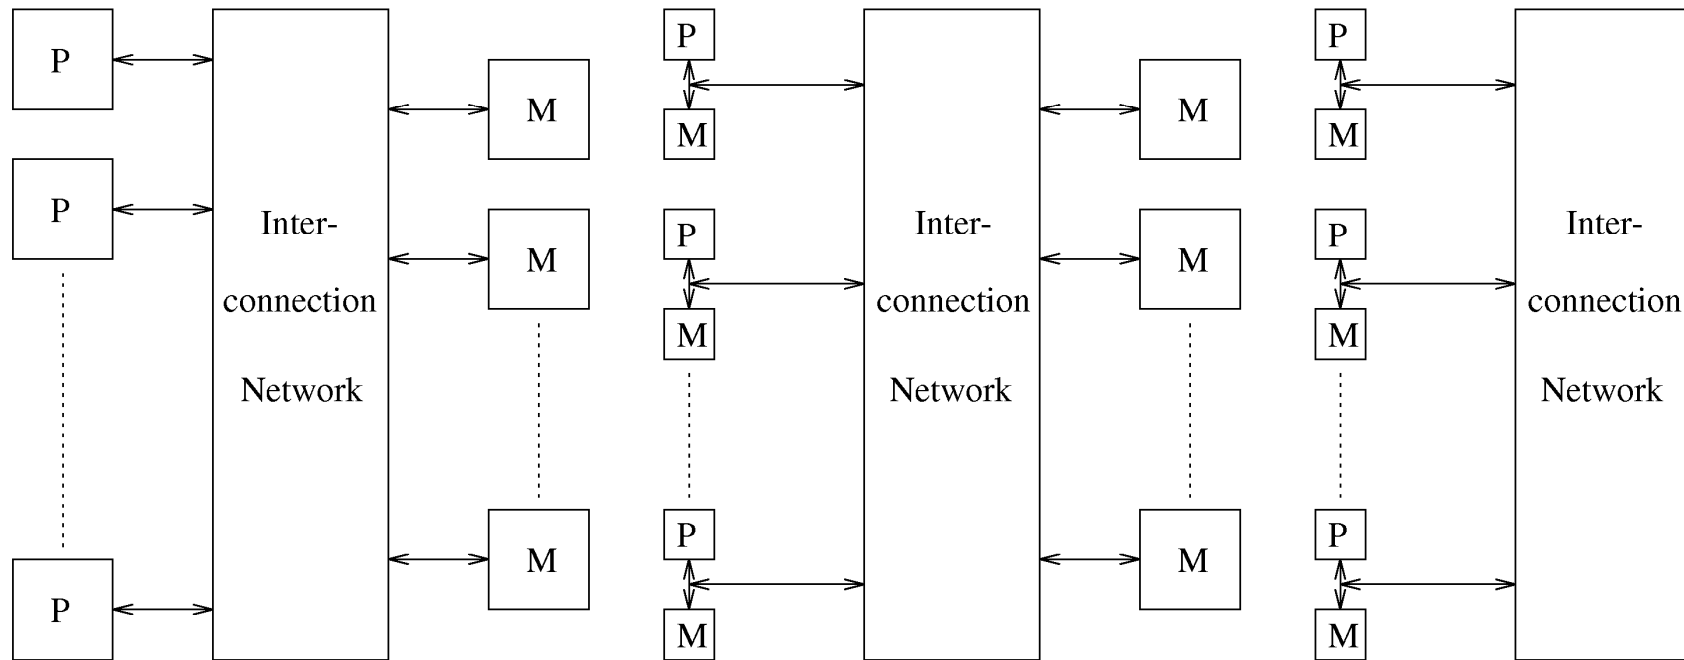
# Parallel Platforms based on Address Space Organization

# Shared Address Space

- Aka shared memory system
- Shared address space:
  - Processors can directly access all the data in the system
  - Inter-processor interaction ?
- Example: multi-core processors



# Shared Address Space




(a)  
Uniform Memory Access (UMA)

(b) (c)  
NonUniform Memory Access (NUMA)

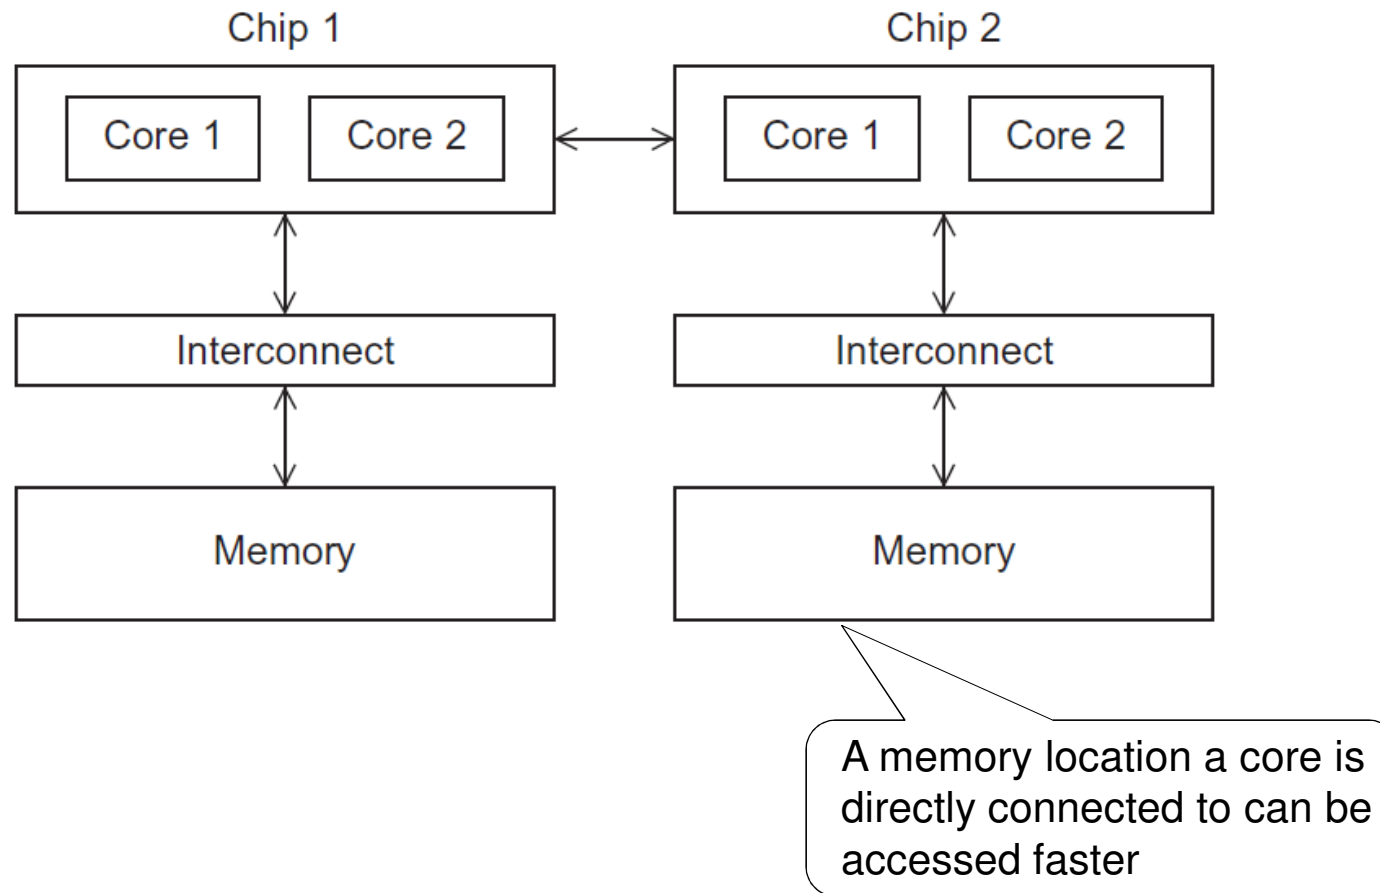


# UMA Multicore System



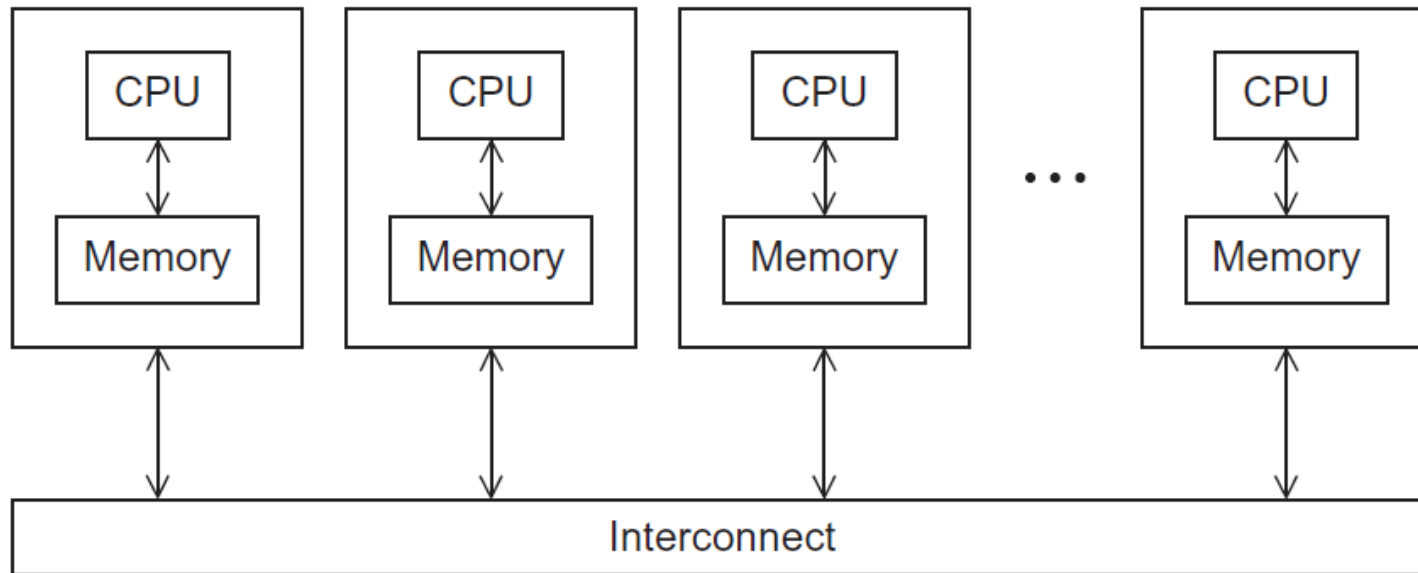
Time to access all the shared memory locations are the same for all the cores

# NUMA Multicore System




# Distributed Address Space

- Aka distributed memory system
- Distributed address space:
  - “Shared nothing:” each processor has a private memory
  - Processors can directly access only local data
  - Inter-processor interaction ?



# Distributed Address Space



# Clusters

- A type of distributed address space machines
- A collection of commodity systems
- Connected by a commodity interconnection network
- Nodes of a cluster are individual computation units joined by a communication network

# Shared vs. Distributed Address

- Shared address:
  - Pro. Vs Con.?
- Distributed address:
  - Pro. Vs. Con.?