


Parallel Programming Models

Parallel Programming Models

- **Programming model** is made up of the languages and libraries that create an abstract view of the machine
- Control
 - How is parallelism created?
 - What orderings exist between operations?
- Data
 - What data is private vs. shared?
 - How is logically shared data accessed or communicated?
- Synchronization
 - What operations can be used to coordinate parallelism?
 - What are the atomic (indivisible) operations?
- Cost
 - How do we account for the cost of each of the above?

Programming Model 1: Shared Memory


- Program is a collection of **threads** of control
 - Can be created dynamically, mid-execution, in some languages
- Each thread has a set of private variables, e.g., local stack variables
- Also a set of shared variables, e.g., static variables, shared common blocks, or global heap
 - Threads communicate implicitly by writing and reading shared variables.
 - Threads coordinate by synchronizing on shared variables



A Simple Example

- Shared memory strategy:
 - small number $p \ll n = \text{size}(A)$ processors
 - attached to single memory
- Parallel Decomposition:
 - Each evaluation and each partial sum is a task.
- Assign n/p numbers to each proc
 - Each computes independent “private” results and partial sum.
 - Collect the p partial sums and compute a global sum.

$$\sum_{i=0}^{n-1} f(A[i])$$



Two Classes of Data:

- Which are logically shared data?
- Which are logically private data?

Shared Memory “Code”

```
fork(sum, a[0:n/2-1]);  
sum(a[n/2, n-1]);
```

```
static int s = 0;
```

Thread 1

```
for i = 0, n/2-1  
  s = s + f(A[i])
```

Thread 2

```
for i = n/2, n-1  
  s = s + f(A[i])
```

- What is the problem with this program?

Shared Memory “Code”

```
static int s = 0;
```

Thread 1

```
local_s1 = 0  
for i = 0, n/2-1
```


Thread 2

```
local_s2 = 0  
for i = n/2, n-1
```

- How to solve the race condition?


Machine Model 1a: Shared Memory

- Processors all connected to a large shared memory
 - Typically called Symmetric Multiprocessors (SMPs)
 - Multicore chips, except that all caches are shared
 - E.g., UMA
- Difficulty scaling to large numbers of processors
 - ≤ 32 processors typical
- Cost: much cheaper to access data in cache than main memory.




Machine Model 1b: Multithreaded Processor

- Multiple thread “contexts” without full processors
- Memory and some other state is shared
- Sun Niagara processor
 - Up to 64 threads all running simultaneously (8 threads x 8 cores)
 - In addition to sharing memory, they share floating point units
- Cray MTA and Eldorado processors (for HPC)



Machine Model 1c: Distributed Shared Memory

- Memory is logically shared, but physically distributed
 - Any processor can access any address in memory
 - Cache lines (or pages) are passed around machine
 - E.g., NUMA
- SGI is canonical example (+ research machines)
 - Scales to 512 (SGI Altix (Columbia) at NASA/Ames)
 - Limitation is *cache coherency protocols*



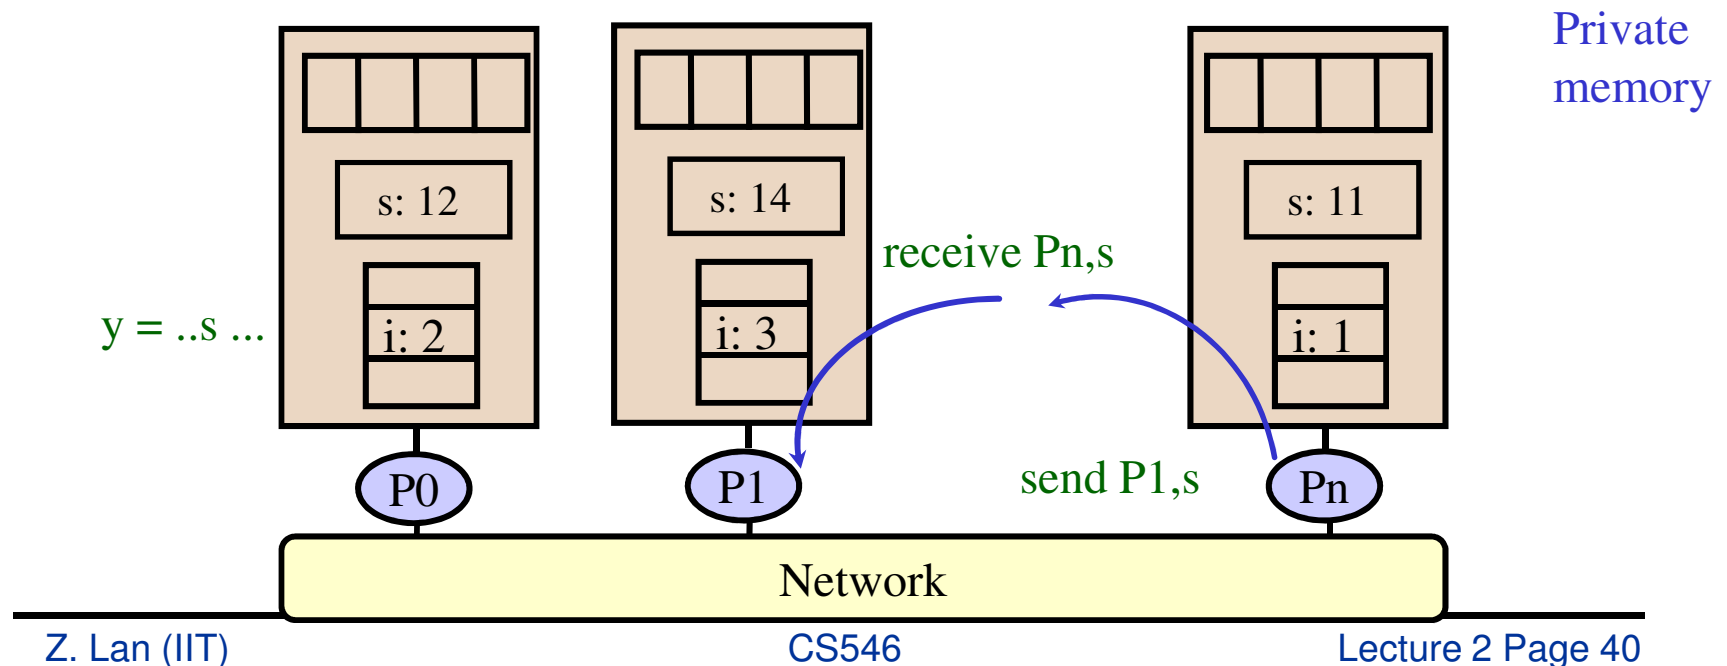
Cache lines (pages) must be large to amortize overhead



locality still critical to performance


Programming Model 2: Message Passing

- Program consists of a collection of named processes
 - Usually fixed at program startup time
 - Thread of control plus local address space -- **NO shared data**
 - Logically shared data is partitioned over local processes.
- Processes communicate by explicit send/receive pairs
 - Coordination is implicit in every communication event



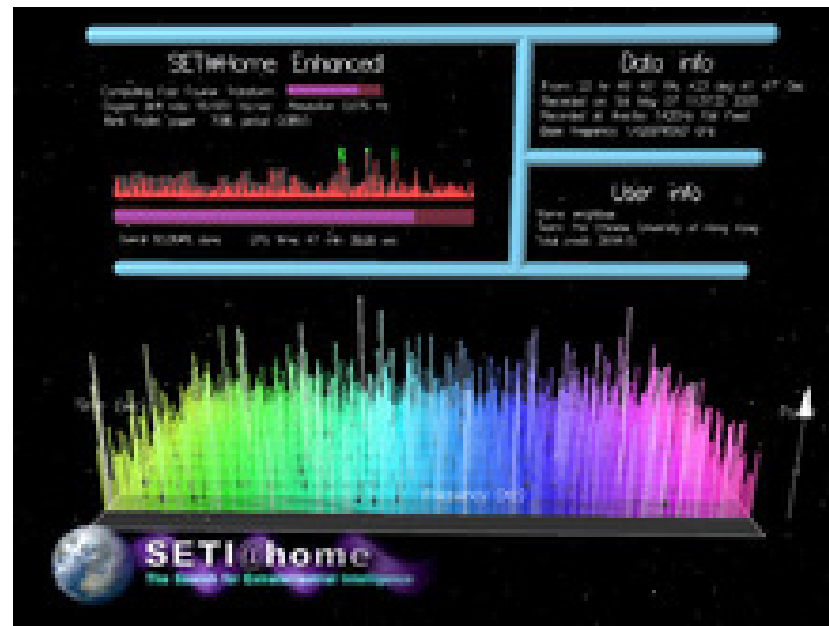
Machine Model 2a: Distributed Memory

- Cray XT4, XT 5
- PC Clusters (Berkeley NOW, Beowulf)
- IBM SP-3, Millennium, CITRIS are distributed memory machines, but the nodes are SMPs.
- Each processor has its own memory and cache but cannot directly access another processor's memory
 - Aka “distributed address space machines”
- Each “node” has a Network Interface (NI) for all communication and synchronization.



Machine Model 2b: Internet/Grids


- SETI@Home: Running on 500,000 PCs
 - ~1000 CPU Years per Day
 - 485,821 CPU Years so far
- Sophisticated Data & Signal Processing Analysis
- Distributes Datasets from Arecibo Radio Telescope



Programming Model 3: Data Parallel


- Single thread of control consisting of parallel operations.
- Parallel operations applied to all (or a defined subset) of a data structure, usually an array
 - Communication is implicit in parallel operators
 - Elegant and easy to understand and reason about
 - Coordination is implicit – statements executed synchronously
- Drawbacks:
 - ???

A = array of all data
fA = f(A)
s = sum(fA)



Machine Model 3a: SIMD System

- A large number of (usually) small processors.
 - A single “control processor” issues each instruction.
 - Each processor executes the same instruction.
 - Some processors may be turned off on some instructions.
- Originally machines were specialized to scientific computing, few made (e.g., CM2)
- Programming model can be implemented in the compiler



Machine Model 3b: Vector Machines

- Vector architectures are based on a single processor
 - Multiple functional units
 - All performing the same operation
 - Instructions may specify large amounts of parallelism (e.g., 64-way) but hardware executes only a subset in parallel
- Historically important
 - Overtaken by MPPs in the 90s
- Re-emerging in recent years
 - At a large scale in the Earth Simulator (NEC SX6) and Cray X1
 - At a small scale in SIMD media extensions to microprocessors
 - SSE, SSE2 (Intel: Pentium/IA64)
 - At a larger scale in GPUs
- Key idea: Compiler does some of the difficult work of finding parallelism, so the hardware doesn't have to

Programming Model 4: Hybrids

- These programming models can be mixed
 - Message passing (MPI) at the top level with shared memory within a node is common
 - New DARPA HPCS languages mix data parallel and threads in a global address space
 - Global address space models can (often) call message passing libraries or vice versa
 - Global address space models can be used in a hybrid mode
 - Shared memory when it exists in hardware
 - Communication (done by the runtime system) otherwise
- For better or worse....

Machine Model 4: Hybrid machines

- Multicore/SMPs are a building block for a larger machine with a network
- Common names:
 - CLUMP = Cluster of SMPs
- Many modern machines look like this:
 - Millennium, IBM SPs, NERSC Franklin, Hopper
- What is an appropriate programming model #4 ???
 - Treat machine as “flat”, always use message passing, even within SMP (simple, but ignores an important part of memory hierarchy).
 - Shared memory within one SMP, but message passing outside of an SMP.
- Graphics or game processors may also be building block

Together...

- Parallel programming models
 - Shared memory (Pthreads and openMP)
 - Message passing (MPI)
 - Massive data parallelism (GPU programming)
 - Hybrid
- Distributed programming models
 - Cloud programming, MapReduce

Data Dependence

- Data dependence
 - True dependence
 - Anti-dependence
 - Output dependence
 - Loop-carried dependence

When can 2 statements execute in parallel?

- On one processor:

statement 1;
statement 2;

- On two processors:

processor1:
statement1;

processor2:
statement2;

When can 2 statements execute in parallel?

- **Possibility 1**

Processor1:
statement1;

Processor2:
statement2;

time
↓

- **Possibility 2**

Processor1:
statement1;

Processor2:
statement2;

time
↓

When can 2 statements execute in parallel?

- Their order of execution must not matter!
- In other words,
 statement1; statement2;
 must be equivalent to
 statement2; statement1;

Example 1

a = 1;

b = 2;

Example 2

```
a = 1;  
b = a;
```


Example 3

$a = f(x);$

$b = a;$

Example 4

```
b = a;  
a = 1;
```

Example 5

a = 1;

a = 2;

True dependence

Statements S1, S2

S2 has a **true dependence** on S1

iff

?

Anti-dependence

Statements S1, S2.

S2 has an **anti-dependence** on S1

iff

?

Output Dependence

Statements S1, S2.

S2 has an **output dependence** on S1

iff

?

When can 2 statements execute in parallel?

S1 and S2 can execute in parallel

iff

there are **no dependences** between S1 and S2

- true dependences
- anti-dependences
- output dependences

Some dependences can be removed.

Example 6

- Most parallelism occurs in loops.

```
for(i=0; i<100; i++)  
    a[i] = i;
```


Example 7

```
for(i=0; i<100; i++) {  
    a[i] = i;  
    b[i] = 2*i;  
}
```

Example 8

```
for(i=0;i<100;i++) a[i] = i;  
for(i=0;i<100;i++) b[i] = 2*i;
```

Example 9

```
for(i=0; i<100; i++)  
    a[i] = a[i] + 100;
```

Example 10

```
for( i=1; i<100; i++ )  
    a[i] = f(a[i-1]);
```

Loop-carried dependence

- A **loop-carried** dependence is a dependence that is present only if the statements are part of the execution of a loop.
- Otherwise, we call it a **loop-independent** dependence.
- Loop-carried dependences prevent loop iteration parallelization.

Example 11

```
for(i=0; i<100; i++ )  
    for(j=1; j<100; j++ )  
        a[i][j] = f(a[i][j-1]);
```

Example 12

```
for( j=1; j<100; j++ )  
    for( i=0; i<100; i++ )  
        a[i][j] = f(a[i][j-1]);
```

Level of loop-carried dependence

- Is the nesting depth of the loop that carries the dependence.
- Indicates which loops can be parallelized.

Be careful ... Example 13

```
printf("a");  
printf("b");
```

Be careful ... Example 14

$a = f(x);$

$b = g(x);$

Be careful ... Example 15

```
for(i=0; i<100; i++)  
    a[i+10] = f(a[i]);
```

Be careful ... Example 16

```
for( i=1; i<100;i++ ) {  
    a[i] = ...;  
    ... = a[i-1];  
}
```

Be careful ... Example 17

```
for( i=0; i<100; i++ )  
    a[i] = f(a[indexa[i]]);
```

An aside

- Parallelizing compilers analyze program dependences to decide parallelization.
- In parallelization by hand, user does the same analysis.
- Compiler more convenient and more correct
- User more powerful, can analyze more patterns.

To remember

- Statement order must not matter.
- Statements must not have dependences.
- Some dependences can be removed.
- Some dependences may not be obvious.

Summary

1. Parallel platforms (~hardware) and programming models (~software)
 - Note: Parallel machine may or may not be tightly coupled to programming model
2. Data dependence
 - Reading: Kumar – ch 1; Foster – ch 1