

# Shared Memory Programming - OpenMP

*Slides from Rice Univ.*

# Topics for Today

---

- Introduction to OpenMP
- OpenMP directives
  - concurrency directives
    - parallel regions
    - loops, sections, tasks
  - synchronization directives
    - reductions, barrier, critical, ordered
  - data handling clauses
    - shared, private, firstprivate, lastprivate
  - tasks
- Performance tuning hints
- Library primitives
- Environment variables

# What is OpenMP?


---

**Open** specifications for **M**ulti **P**rocessing

- An API for explicit multi-threaded, shared memory parallelism
- Three components
  - compiler directives
  - runtime library routines
  - environment variables
- Higher-level programming model than Pthreads
  - implicit mapping and load balancing of work
- Portable
  - API is specified for C/C++ and Fortran
  - implementations on almost all platforms
- Standardized

# OpenMP at a Glance

---



# OpenMP Is Not

---

- An automatic parallel programming model
  - parallelism is explicit
  - programmer full control (and responsibility) over parallelization
- Meant for distributed-memory parallel systems (by itself)
  - designed for shared address spaced machines
- Necessarily implemented identically by all vendors
- Guaranteed to make the most efficient use of shared memory
  - no data locality control


# OpenMP Targets Ease of Use

---

- OpenMP does not require that single-threaded code be changed for threading
  - enables incremental parallelization of a serial program
- OpenMP only adds compiler directives
  - pragmas (C/C++); significant comments in Fortran
    - if a compiler does not recognize a directive, it simply ignores it
  - simple & limited set of directives for shared memory programs
  - significant parallelism possible using just 3 or 4 directives
    - both coarse-grain and fine-grain parallelism
- If OpenMP is disabled when compiling a program, the program will execute sequentially

# OpenMP: Fork-Join Parallelism

- OpenMP program begins execution as a single master thread
- Master thread executes sequentially until 1<sup>st</sup> parallel region
- When a parallel region is encountered, master thread
  - creates a group of threads
  - becomes the master of this group of threads
  - is assigned the thread id 0 within the group



# OpenMP Directive Format

---

- OpenMP directive forms
  - C and C++ use compiler directives
    - prefix: **#pragma ...**
  - Fortran uses significant comments
    - prefixes: **!\$omp, c\$omp, \*\$omp**
- A directive consists of a directive name followed by clauses

C: **#pragma omp parallel default(shared) private(beta,pi)**  
Fortran: **!\$omp parallel default(shared) private(beta,pi)**



# OpenMP **parallel** Region Directive

---

**#pragma omp parallel** [clause list]

## Typical clauses in [clause list]

A few more clauses  
on slide 37

- Conditional parallelization
  - **if** (scalar expression)
    - determines whether the **parallel** construct creates threads
- Degree of concurrency
  - **num\_threads** (integer expression) : # of threads to create
- Data Scoping
  - **private** (variable list)
    - specifies variables local to each thread
  - **firstprivate** (variable list)
    - similar to the private
    - private variables are initialized to variable value before the parallel directive
  - **shared** (variable list)
    - specifies that variables are shared across all the threads
  - **default** (data scoping specifier)
    - default data scoping specifier may be **shared** or **none**

# Interpreting an OpenMP Parallel Directive

---

```
#pragma omp parallel if (is_parallel==1) num_threads(8) \  
    shared (b) private (a) firstprivate(c) default(none)  
{  
    /* structured block */  
}
```

## Meaning

- **if (is\_parallel== 1) num\_threads(8)**  
—If the value of the variable `is_parallel` is one, create 8 threads
- **shared (b)**  
—each thread shares a single copy of variable `b`
- **private (a) firstprivate(c)**  
—each thread gets private copies of variables `a` and `c`  
—each private copy of `c` is initialized with the value of `c` in main thread when the parallel directive is encountered
- **default(none)**  
— default state of a variable is specified as **none** (rather than **shared**)  
—signals error if not all variables are specified as **shared** or **private**

# Meaning of OpenMP Parallel Directive

```
int a, b;
main() {
    [ // serial segment
      #pragma omp parallel num_threads (8) private (a) shared (b)
      { [ // parallel segment
        ]
      }
    [ // rest of serial segment
    ]
}
```

Sample OpenMP program

**OpenMP**

```
int a, b;
main() {
    [ // serial segment
    ]
    Code
    inserted
    by
    the OpenMP
    compiler [
        for (i = 0; i < 8; i++)
            pthread_create (....., internal_thread_fn_name, ...);
        for (i = 0; i < 8; i++)
            pthread_join (.....);
    ]
    [ // rest of serial segment
    ]
}

void *internal_thread_fn_name (void *packaged_argument) {
    int a;
    [ // parallel segment
    ]
}
```

Corresponding Pthreads translation

**Pthreads  
equivalent**

# Specifying Worksharing

---

**Within the scope of a parallel directive, worksharing directives allow concurrency between iterations or tasks**

- **OpenMP provides two directives**
  - **DO/for**: concurrent loop iterations
  - **sections**: concurrent tasks

# Worksharing **DO/for** Directive

---

**for** directive partitions parallel iterations across threads

**DO** is the analogous directive for Fortran

- Usage:

```
#pragma omp for [clause list]
```

```
/* for loop */
```

- Possible clauses in [clause list]

- **private**, **firstprivate**, **lastprivate**

- **reduction**

- **schedule**, **nowait**, **and ordered**

- Implicit barrier at end of **for** loop

# A Simple Example Using **parallel** and **for**

---

## Program

```
void main()  {  
#pragma omp parallel num_threads(3)  
{  
    int i;  
    printf("Hello world\n");  
    #pragma omp for  
    for (i = 1; i <= 4; i++) {  
        printf("Iteration %d\n", i);  
    }  
    printf("Goodbye world\n");  
}  
}
```

## Output

```
Hello world  
Hello world  
Hello world  
Iteration 1  
Iteration 2  
Iteration 3  
Iteration 4  
Goodbye world  
Goodbye world  
Goodbye world
```

# Reduction Clause for Parallel Directive

---

Specifies how to combine local copies of a variable in different threads into a single copy at the master when threads exit

- Usage: **reduction** (operator: variable list)  
—variables in list are implicitly private to threads
- Reduction operators: **+**, **\***, **-**, **&**, **|**, **^**, **&&**, and **||**
- Usage sketch

```
#pragma omp parallel reduction(+: sum) num_threads(8)
{
/* compute local sum in each thread here */
}
/* sum here contains sum of all local instances of sum */
```

# OpenMP Reduction Clause Example

---

## OpenMP threaded program to estimate PI

```
#pragma omp parallel default(private) shared (npoints) \  
    reduction(+: sum) num_threads(8)
```

```
{
```

```
    num_threads = omp_get_num_threads();  
    sample_points_per_thread = npoints / num_threads;  
    sum = 0;
```

```
    for (i = 0; i < sample_points_per_thread; i++) {  
        coord_x = (double)(rand_r(&seed)) / ((double)((2<<14)-1) - 0.5);  
        coord_y = (double)(rand_r(&seed)) / ((double)((2<<14)-1) - 0.5);  
        if ((coord_x * coord_x + coord_y * coord_y) < 0.25)  
            sum ++;
```

```
    }
```

```
}
```

here, user  
manually  
divides work

- a local copy of sum for each thread
- all local copies of sum added together and stored in master



# Using Worksharing **for** Directive

---

```
#pragma omp parallel default(private) shared (npoints) \  
    reduction(+: sum) num_threads(8)  
{  
    sum = 0;  
    #pragma omp for  
    for (i = 0; i < npoints; i++) {  
        rand_no_x = (double)(rand_r(&seed)) / (double)((2<<14)-1);  
        rand_no_y = (double)(rand_r(&seed)) / (double)((2<<14)-1);  
        if (((rand_no_x - 0.5) * (rand_no_x - 0.5) +  
            (rand_no_y - 0.5) * (rand_no_y - 0.5)) < 0.25)  
            sum ++;  
    }  
}
```

worksharing **for**  
divides work

Implicit barrier at end of loop

# Mapping Iterations to Threads

---

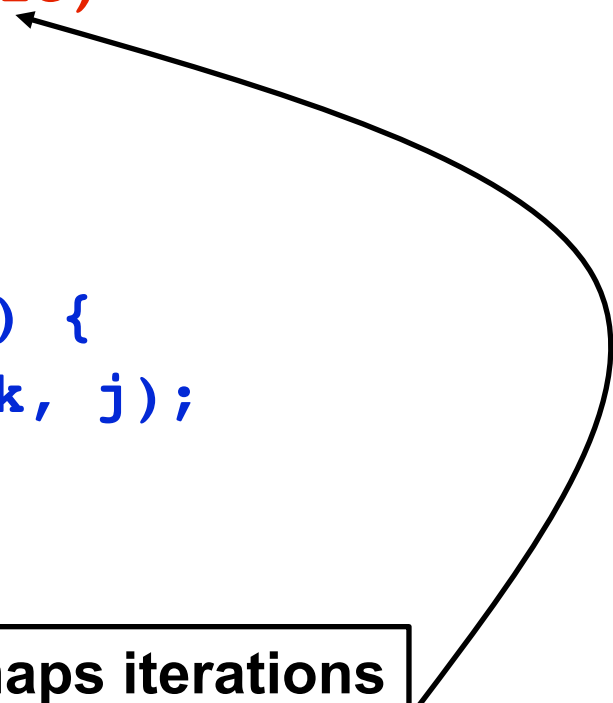
**schedule** clause of the **for** directive

- Recipe for mapping iterations to threads
- Usage: **schedule**(**scheduling\_class**[ , **parameter**] ).
- Four scheduling classes
  - **static**: work partitioned at compile time
    - iterations statically divided into pieces of size *chunk*
    - statically assigned to threads
  - **dynamic**: work evenly partitioned at run time
    - iterations are divided into pieces of size *chunk*
    - chunks dynamically scheduled among the threads
    - when a thread finishes one chunk, it is dynamically assigned another
    - default chunk size is 1
  - **guided**: guided self-scheduling
    - chunk size is exponentially reduced with each dispatched piece of work
    - the default minimum chunk size is 1
  - **runtime**:
    - scheduling decision from environment variable **OMP\_SCHEDULE**
    - illegal to specify a chunk size for this clause.

# Statically Mapping Iterations to Threads

---

```
/* static scheduling of matrix multiplication loops */
#pragma omp parallel default(private) \
    shared (a, b, c, dim) num_threads(4)
#pragma omp for schedule(static)
for (i = 0; i < dim; i++) {
    for (j = 0; j < dim; j++) {
        c(i,j) = 0;
        for (k = 0; k < dim; k++) {
            c(i,j) += a(i, k) * b(k, j);
        }
    }
}
```



**static schedule maps iterations  
to threads at compile time**

# Avoiding Unwanted Synchronization

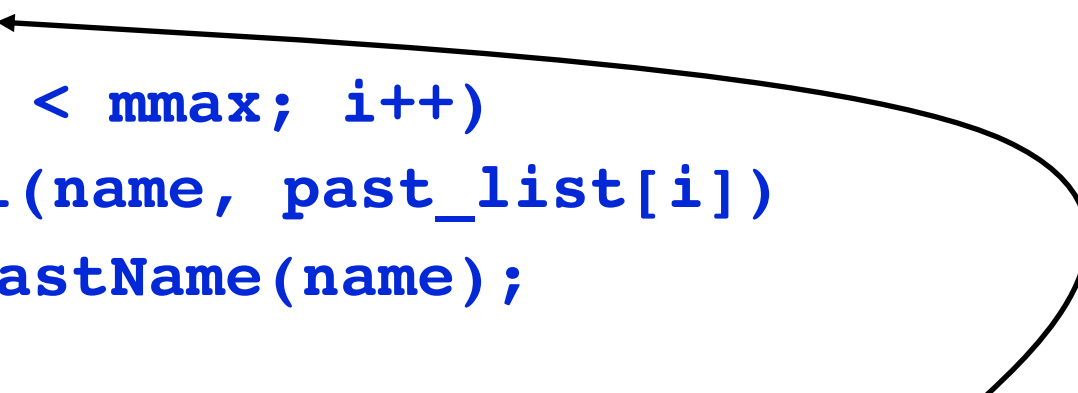
---

- Default: worksharing **for** loops end with an implicit barrier
- Often, less synchronization is appropriate
  - series of independent **for**-directives within a parallel construct
- **nowait** clause
  - modifies a **for** directive
  - avoids implicit barrier at end of for

# Avoiding Synchronization with **nowait**

---

```
#pragma omp parallel
{
    #pragma omp for nowait
        for (i = 0; i < nmax; i++)
            if (isEqual(name, current_list[i])
                processCurrentName(name);
    #pragma omp for
        for (i = 0; i < mmax; i++)
            if (isEqual(name, past_list[i])
                processPastName(name);
}
```



any thread can begin second loop immediately without waiting for other threads to finish first loop

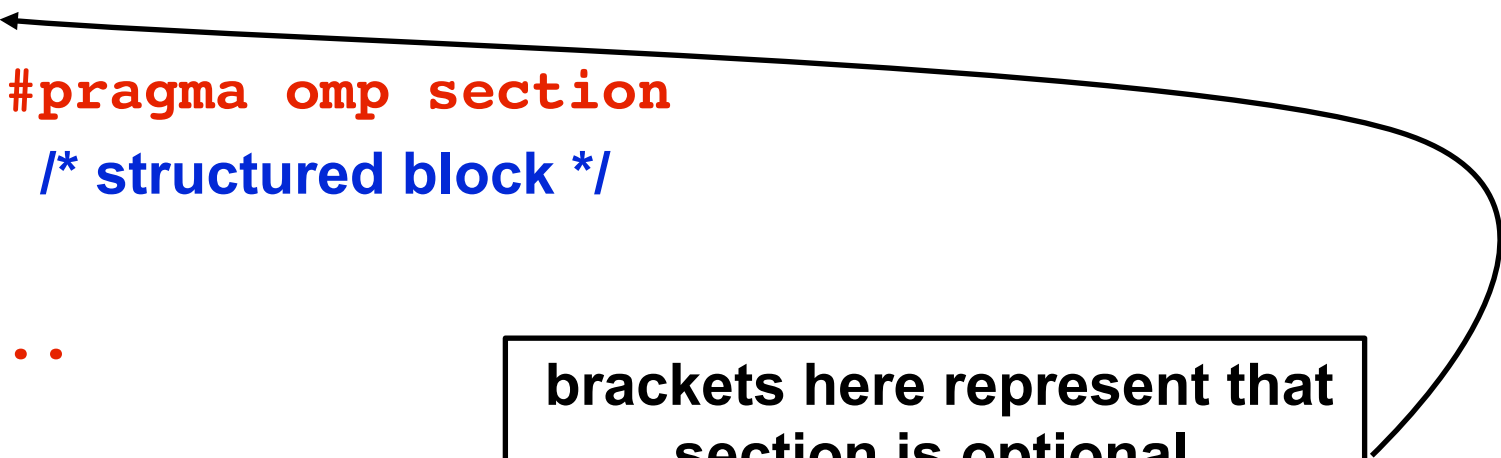
# Worksharing **sections** Directive

---

**sections** directive enables specification of task parallelism

- Usage

```
#pragma omp sections [clause list]
{
    [#pragma omp section
    /* structured block */
]
    [#pragma omp section
    /* structured block */
]
    ...
}
```



brackets here represent that  
section is optional,  
not the syntax for using them

# Using the **sections** Directive

```
#pragma omp parallel
```

parallel section encloses all parallel work

```
{
```

```
    #pragma omp sections
```

sections: task parallelism

```
{
```

```
    #pragma omp section
```

```
{
```

```
        taskA();
```

```
}
```

```
    #pragma omp section
```

```
{
```

```
        taskB();
```

```
}
```

```
    #pragma omp section
```

```
{
```

```
        taskC();
```

```
}
```

```
}
```


```
}
```

three concurrent tasks  
need not be procedure calls



# Nesting **parallel** Directives

- Nested parallelism enabled using the **OMP\_NESTED** environment variable
  - **OMP\_NESTED = TRUE** → nested parallelism is enabled
- Each parallel directive creates a new team of threads





# Synchronization Constructs in OpenMP

---


**#pragma omp barrier** wait until all threads arrive here

**#pragma omp single [clause list]**

structured block

**#pragma omp master**

structured block



single-threaded  
execution

Use MASTER instead of SINGLE wherever possible

- **MASTER** = IF-statement with no implicit **BARRIER**
  - equivalent to  
`IF(omp_get_thread_num() == 0) {...}`
- **SINGLE**: implemented like other worksharing constructs
  - keeping track of which thread reached SINGLE first adds overhead

# Synchronization Constructs in OpenMP

---


**#pragma omp critical [(name)]** critical section: like a named lock  
structured block

**#pragma omp ordered** for loops with carried dependences  
structured block

# Example Using **critical**

---

```
#pragma omp parallel
{
#pragma omp for nowait shared(best_cost)
  for (i = 0; i < nmax; i++) {
    int my_cost;
    ...
#pragma omp critical
    {
      if (best_cost < my_cost)
        best_cost = my_cost;
    }
    ...
  }
}
```



**critical ensures mutual exclusion  
when accessing shared state**

# Example Using **ordered**

---

```
#pragma omp parallel
{
#pragma omp for nowait shared(a)
  for (k = 0; k < nmax; k++) {
    ...
#pragma omp ordered
  {
    a[k] = a[k-1] + ...;
  }
  ...
}
```



ordered ensures carried dependence  
does not cause a data race

# Orphaned Directives

- Directives may not be lexically nested in a parallel region
  - may occur in a separate program unit

```
...  
!$omp parallel  
call phase1  
call phase2  
!$omp end parallel  
...
```

```
subroutine phase1  
!$omp do private(i) shared(n)  
do i = 1, n  
call some_work(i)  
end do  
!$omp end do  
end
```

```
subroutine phase2  
!$omp do private(j) shared(n)  
do j = 1, n  
call more_work(j)  
end do  
!$omp end do  
end
```

- Dynamically bind to enclosing parallel region at run time
- Benefits
  - enables parallelism to be added with a minimum of restructuring
  - improves performance: enables single parallel region to bind with worksharing constructs in multiple called routines
- Execution rules
  - orphaned worksharing construct is executed serially when not called from within a parallel region

# OpenMP 3.0 Tasks

---

- **Motivation: support parallelization of irregular problems**
  - unbounded loops
  - recursive algorithms
  - producer consumer
- **What is a task?**
  - work unit
    - execution can begin immediately, or be deferred
  - components of a task
    - code to execute, data environment, internal control variables
- **Task execution**
  - data environment is constructed at creation
  - tasks are executed by threads of a team
  - a task can be tied to a thread (i.e. migration/stealing not allowed)
    - by default: a task is tied to the first thread that executes it

# OpenMP 3.0 Tasks

---

**#pragma omp task** [clause list]

## Possible clauses in [clause list]

- Conditional parallelization
  - **if** (scalar expression)
    - determines whether the construct creates a task
- Binding to threads
  - **untied**
- Data scoping
  - **private** (variable list)
    - specifies variables local to the child task
  - **firstprivate** (variable list)
    - similar to the private
    - private variables are initialized to value in parent task before the directive
  - **shared** (variable list)
    - specifies that variables are shared with the parent task
  - **default** (data handling specifier)
    - default data handling specifier may be **shared** or **none**

# Composing Tasks and Regions

---

```
#pragma omp parallel
{
    #pragma omp task
        x();
    #pragma omp barrier
    #pragma omp single
    {
        #pragma omp task
            y();
    }
}
```

one  $\underline{x}$  task created for each thread in the parallel region

all  $\underline{x}$  tasks complete at barrier

one  $\underline{y}$  task created

region end:  $\underline{y}$  task completes



# Data Scoping for Tasks is Tricky

---

**If no default clause specified**

- Static and global variables are **shared**
- Automatic (local) variables are **private**
- Variables for orphaned tasks are **firstprivate** by default
- Variables for non-orphaned tasks inherit the shared attribute
  - task variables are **firstprivate** unless **shared** in the enclosing context

# Fibonacci using (Orphaned) OpenMP 3.0 Tasks

```
int fib ( int n )
{
    int x,y;
    if ( n < 2 ) return n;
    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);
    #pragma omp taskwait
    return x + y;
}
```

need **shared** for x and y;  
default would be  
**firstprivate**

suspend parent task until  
children finish

```
int main (int argc, char **argv)
{
    int n, result;
    n = atoi(argv[1]);
    #pragma omp parallel
    {
        #pragma omp single
        {
            result = fib(n);
        }
    }

    printf("fib(%d) = %d\n",
        n, result);
}
```

create team  
of threads to  
execute tasks

only one thread  
performs the  
outermost call

# List Traversal

---

```
Element first, e;  
#pragma omp parallel  
#pragma omp single  
{  
    for (e = first; e; e = e->next)  
#pragma omp task firstprivate(e)  
        process(e);  
}
```

Is the use of variables safe as written?

# Task Scheduling

---

- **Tied tasks**

- only the thread that the task is tied to may execute it
- task can only be suspended at a suspend point
  - task creation
  - task finish
  - taskwait
  - barrier
- if a task is not suspended at a barrier, it can only switch to a descendant of any task tied to the thread

- **Untied tasks**

- no scheduling restrictions
  - can suspend at any point
  - can switch to any task
- implementation may schedule for locality and/or load balance

# Summary of Clause Applicability

Clause	Directive					
	PARALLEL	DO/for	SECTIONS	SINGLE	PARALLEL DO/for	PARALLEL SECTIONS
IF	●				●	●
PRIVATE	●	●	●	●	●	●
SHARED	●	●			●	●
DEFAULT	●				●	●
FIRSTPRIVATE	●	●	●	●	●	●
LASTPRIVATE		●	●		●	●
REDUCTION	●	●	●		●	●
COPYIN	●				●	●
SCHEDULE		●			●	
ORDERED		●			●	
NOWAIT		●	●	●		

# Performance Tuning Hints

---

Parallelize at the highest level, e.g. outermost **DO**/**for** loops

```
!$OMP PARALLEL
....
do j = 1, 20000
!$OMP DO
  do k = 1, 10000
    ...
  enddo !k
!$OMP END DO
enddo !j
...
!$OMP END PARALLEL
```

Slower

```
!$OMP PARALLEL
....
!$OMP DO
do k = 1, 10000
  do j = 1, 20000
    ...
  enddo !j
enddo !k
!$OMP END DO
...
!$OMP END PARALLEL
```

Faster

# Performance Tuning Hints

---

Merge independent parallel loops when possible

```
!$OMP PARALLEL DO
....
!$OMP DO
  statement 1
!$OMP END DO
!$OMP DO
  statement 2
!$OMP END DO
....
!$OMP END PARALLEL
```

Slower

```
!$OMP PARALLEL DO
....
!$OMP DO
  statement 1
  statement 2
!$OMP END DO
....
!$OMP END PARALLEL
```

Faster

# Performance Tuning Hints

---

## Minimize use of synchronization

- **BARRIER**
- **CRITICAL** sections
  - if necessary, use named **CRITICAL** for fine-grained locking
- **ORDERED** regions
- Use **NOWAIT** clause to avoid unnecessary barriers
  - adding **NOWAIT** to a region's final DO eliminates a redundant barrier
- Use explicit **FLUSH** with care
  - flushes can evict cached values
  - subsequent data accesses may require reloads from memory



# OpenMP Library Functions

---

- **Processor count**

```
int omp_get_num_procs(); /* # PE currently available */
```

```
int omp_in_parallel(); /* determine whether running in parallel */
```

- **Thread count and identity**

```
/* max # threads for next parallel region. only call in serial region */
```

```
void omp_set_num_threads(int num_threads);
```

```
int omp_get_num_threads(); /*# threads currently active */
```

# OpenMP Library Functions

---

- **Controlling and monitoring thread creation**

```
void omp_set_dynamic (int dynamic_threads);  
int omp_get_dynamic ();  
void omp_set_nested (int nested);  
int omp_get_nested ();
```

- **Mutual exclusion**

```
void omp_init_lock(omp_lock_t *lock);  
void omp_destroy_lock(omp_lock_t *lock);  
  
void omp_set_lock(omp_lock_t *lock);  
void omp_unset_lock(omp_lock_t *lock);  
int omp_test_lock(omp_lock_t *lock);
```

—Lock routines have a nested lock counterpart for recursive mutexes

# OpenMP Environment Variables

---

- **OMP\_NUM\_THREADS**  
— specifies the default number of threads for a parallel region
- **OMP\_DYNAMIC**  
— specifies if the number of threads can be dynamically changed
- **OMP\_NESTED**  
— enables nested parallelism (may be nominal: one thread)
- **OMP\_SCHEDULE**  
— specifies scheduling of **for**-loops if the clause specifies runtime
- **OMP\_STACKSIZE** (for non-master threads)
- **OMP\_WAIT\_POLICY** (ACTIVE or PASSIVE)
- **OMP\_MAX\_ACTIVE\_LEVELS**  
— integer value for maximum # nested parallel regions
- **OMP\_THREAD\_LIMIT** (# threads for entire program)

OpenMP 3.0

# OpenMP Directives vs. Pthreads

---

- **Directive advantages**
  - directives facilitate a variety of thread-related tasks
  - frees programmer from
    - initializing attribute objects
    - setting up thread arguments
    - partitioning iteration spaces, ...
- **Directive disadvantages**
  - data exchange is less apparent
    - leads to mysterious overheads  
data movement, false sharing, and contention
  - API is less expressive than Pthreads
    - lacks condition waits, locks of different types, and flexibility for building composite synchronization operations

# The Future of OpenMP

---

- OpenMP 3.1 standard finalized, July 2011
- OpenMP 4.0 standardization process has begun; topics under discussion are the following:
  - support for heterogeneous systems (GPUs etc.)
  - tasking model refinement
  - locality and affinity
  - thread team control
  - transactional memory and thread-level speculation
  - additional synchronization mechanisms
  - OpenMP error model
  - interoperability and composability
  - tools support in Spec
  - consideration of Fortran 2003, other language bindings (Java, Python)